

# Reinforcement Learning

## Monte Carlo Methods

---

Stefano Albrecht, Pavlos Andreadis

28 January 2020



THE UNIVERSITY of EDINBURGH  
**informatics**


- Monte Carlo policy evaluation
- Monte Carlo control with...
  - Exploring starts
  - Soft policies
  - Off-policy learning

## Recap: Generalised Policy Iteration

DP methods iterate through policy evaluation and improvement until convergence to optimal value function  $v_*$  and policy  $\pi_*$

- Policy evaluation via repeated application of Bellman operator
- Requires **complete knowledge** of MDP model:  $p(s', r|s, a)$

*Can we compute optimal policy without knowledge of complete model?*



# Monte Carlo Policy Evaluation

**Monte Carlo (MC)** methods **learn** value function based on **experience**

- Experience: entire episodes  $E^i = \langle S_0^i, A_0^i, R_1^i, S_1^i, A_1^i, R_2^i, \dots, S_{T_i}^i \rangle$

Two ways to obtain episodes:

- **Real experience:** generate episodes directly from “real world”
- **Simulated experience:** use simulation model  $\hat{p}$  to sample episodes
  - $\hat{p}(s, a)$  returns a pair  $(s', r)$  with probability  $p(s', r | s, a)$

MC does not require complete model  $p(s', r | s, a)$

# Monte Carlo Policy Evaluation

**Monte Carlo (MC)** methods **learn** value function based on **experience**

- Estimate value function by averaging sample returns:

$$v_{\pi}(s) \doteq \mathbb{E}_{\pi} \left[ \sum_{k=0}^{T-1} \gamma^k R_{k+1} | S_t = s \right] \approx \frac{1}{|\mathcal{E}(s)|} \sum_{t_i \in \mathcal{E}(s)} \sum_{k=t_i}^{T_i-1} \gamma^{k-t_i} R_{k+1}^i$$

where for each past episode  $E^i = \langle S_0^i, A_0^i, R_1^i, S_1^i, A_1^i, R_2^i, \dots, S_{T_i}^i \rangle$ :

- **First-visit MC**:  $\mathcal{E}(s)$  contains *first* time  $t_i$  for which  $S_{t_i}^i = s$  in  $E^i$
- **Every-visit MC**:  $\mathcal{E}(s)$  contains *all* times  $t_i$  for which  $S_{t_i}^i = s$  in  $E^i$

- Both methods converge to  $v_{\pi}(s)$  as  $|\mathcal{E}(s)| \rightarrow \infty$

# First-Visit Monte Carlo Policy Evaluation

Initialize:

$\pi \leftarrow$  policy to be evaluated

$V \leftarrow$  an arbitrary state-value function

$Returns(s) \leftarrow$  an empty list, for all  $s \in \mathcal{S}$

Repeat forever:

Generate an episode using  $\pi$

For each state  $s$  appearing in the episode:

$G \leftarrow$  return following the first occurrence of  $s$

Append  $G$  to  $Returns(s)$

$V(s) \leftarrow \text{average}(Returns(s))$

## Example: Blackjack

Initial state:

Player

*Ace worth 1  
or 11*



Dealer



*Hidden card*

First, player samples  
cards from deck (hit)  
until stop (stick)

Then, dealer samples  
cards from deck (hit)  
until sum  $> 16$  (stick)

Player loses (-1 reward) if bust (card sum  $> 21$ )

Player wins (+1 reward) if Dealer bust or Player sum  $>$  Dealer sum

## Example: Blackjack

### Player policy $\pi$ :

stick if player sum is  
20 or 21, else hit

Estimate of  $v_\pi$  using MC ...

### States (3-tuple):

- Player sum (12–21)
- Dealer card (ace–10)
- Usable ace?



# Example: Blackjack

Player policy  $\pi$ :

stick if player sum is 20 or 21, else hit



States (3-tuple):

- Player sum (12–21)
- Dealer card (ace–10)
- Usable ace?



Usable  
ace

No  
usable  
ace

After 10,000 episodes



After 500,000 episodes



# States in Blackjack

Couldn't we just define states as  $S_t = \{\text{Player cards, Dealer card}\}$ ?

- Tricky: states would have variable length (player cards)
- If we fix maximum number of player cards to 4, then there are  $10^5 = 100,000$  possible states! (ignoring face cards and ordering)

# States in Blackjack

Couldn't we just define states as  $S_t = \{\text{Player cards, Dealer card}\}$ ?

- Tricky: states would have variable length (player cards)
- If we fix maximum number of player cards to 4, then there are  $10^5 = 100,000$  possible states! (ignoring face cards and ordering)

Blackjack example uses **engineered state features**:

- Fixed length:  $S_t = (\text{Player sum, Dealer card, Usable ace?})$
- Player sum limited to range 12–21 because decision below 12 is trivial (always hit)
- Number of states:  $10 * 10 * 2 = 200 \rightarrow$  much smaller problem!
- Still has all relevant information

# Blackjack and Dynamic Programming

Can we solve Blackjack MDP with DP methods?

- Yes, in principle, because we know complete MDP
- But computing  $p(s', r|s, a)$  can be complicated!  
E.g. what is probability of +1 reward as function of Dealer's showing card?
- On other hand, easy to code a simulation model:
  - Use Dealer rule to sample cards until stick/bust, then compute reward
  - Reward outcome is distributed by  $p(s', r|s, a)$
- MC can evaluate policy without knowledge of probabilities  $p(s', r|s, a)$


# Monte Carlo Estimation of Action Values

MC methods can learn  $v_\pi$  without knowledge of model  $p(s', r|s, a)$

⇒ But improving policy  $\pi$  from  $v_\pi$  requires model (*why?*)

Must estimate **action values**:


$$q_\pi(s, a) \doteq \mathbb{E}_\pi[G_t | S_t = s, A_t = a]$$



- Improve policy without model:  $\pi'(s) = \arg \max_a q_\pi(s, a)$
- Use same MC methods to learn  $q_\pi$ , but visits are to  $(s, a)$  pairs
- Converges to  $q_\pi$  if every  $(s, a)$  pair visited infinitely many times in limit

E.g. **exploring starts**: every  $(s, a)$  pair has non-zero probability of being starting pair of episode


- **MC policy evaluation:**  
Estimate  $q_\pi$  using MC method
- **Policy improvement:**  
Improve  $\pi$  by making greedy wrt  $q_\pi$



# Monte Carlo Control with Exploring Starts

Greedy policy meets conditions for policy improvement theorem:

$$\begin{aligned} q_{\pi_k}(s, \pi_{k+1}(s)) &= q_{\pi_k}(s, \arg \max_a q_{\pi_k}(s, a)) \\ &= \max_a q_{\pi_k}(s, a) \\ &\geq q_{\pi_k}(s, \pi_k(s)) \\ &= v_{\pi_k}(s) \end{aligned}$$



Assumes exploring starts and *infinite* MC iterations (*why?*)

- In practice, update only to a given performance threshold
- Or alternate between evaluation and improvement per episode

## Monte Carlo Control with Exploring Starts

Initialize, for all  $s \in \mathcal{S}$ ,  $a \in \mathcal{A}(s)$ :

$Q(s, a) \leftarrow$  arbitrary

$\pi(s) \leftarrow$  arbitrary

$Returns(s, a) \leftarrow$  empty list

Repeat forever:

Choose  $S_0 \in \mathcal{S}$  and  $A_0 \in \mathcal{A}(S_0)$  s.t. all pairs have probability  $> 0$

Generate an episode starting from  $S_0, A_0$ , following  $\pi$

For each pair  $s, a$  appearing in the episode:

$G \leftarrow$  return following the first occurrence of  $s, a$

Append  $G$  to  $Returns(s, a)$

$Q(s, a) \leftarrow \text{average}(Returns(s, a))$

For each  $s$  in the episode:

$\pi(s) \leftarrow \operatorname{argmax}_a Q(s, a)$






# Blackjack Example with MC-ES

Policy  $\pi$ :

stick if player sum  
is 20 or 21, else hit

Exploring starts:

sample initial states  
uniformly randomly



# Monte Carlo Control With Soft Policies

Convergence to  $q_\pi$  requires that all  $(s, a)$ -pairs are visited infinitely many times

- Exploring starts guarantee this, but impractical (*why?*)

Other approach: use **soft policy** such that  $\pi(a|s) > 0$  for all  $s, a$

- e.g.  $\epsilon$ -soft policy:  $\pi(a|s) \geq \epsilon/|\mathcal{A}|$  for  $\epsilon > 0$
- **Policy improvement:** make policy  $\epsilon$ -greedy wrt  $q_\pi$

$$\pi'(a|s) \doteq \begin{cases} \epsilon/|\mathcal{A}| + (1 - \epsilon) & \text{if } a = \arg \max_{a'} q_\pi(s, a') \\ \epsilon/|\mathcal{A}| & \text{else} \end{cases}$$

# Monte Carlo Control With Soft Policies

$\epsilon$ -greedy policy meets conditions for policy improvement theorem:

$$\begin{aligned} q_{\pi}(s, \pi'(s)) &= \sum_a \pi'(a|s) q_{\pi}(s, a) \\ &= \frac{\epsilon}{|\mathcal{A}|} \sum_a q_{\pi}(s, a) + (1 - \epsilon) \max_a q_{\pi}(s, a) \\ &\geq \frac{\epsilon}{|\mathcal{A}|} \sum_a q_{\pi}(s, a) + (1 - \epsilon) \sum_a \frac{\pi(a|s) - \epsilon/|\mathcal{A}|}{1 - \epsilon} q_{\pi}(s, a) \\ &= \frac{\epsilon}{|\mathcal{A}|} \sum_a q_{\pi}(s, a) - \frac{\epsilon}{|\mathcal{A}|} \sum_a q_{\pi}(s, a) + \sum_a \pi(a|s) q_{\pi}(s, a) \\ &= v_{\pi}(s) \end{aligned}$$

- Thus,  $\pi'$  better or equal to  $\pi$ , but both are still  $\epsilon$ -soft
- $q_{\pi}(s, \pi'(s)) = v_{\pi}(s)$  only when  $\pi'$  and  $\pi$  both optimal  $\epsilon$ -soft policies

## Monte Carlo Control With Soft Policies

Initialize, for all  $s \in \mathcal{S}$ ,  $a \in \mathcal{A}(s)$ :

$Q(s, a) \leftarrow$  arbitrary

$Returns(s, a) \leftarrow$  empty list

$\pi(a|s) \leftarrow$  an arbitrary  $\varepsilon$ -soft policy

Repeat forever:

(a) Generate an episode using  $\pi$

(b) For each pair  $s, a$  appearing in the episode:

$G \leftarrow$  return following the first occurrence of  $s, a$

Append  $G$  to  $Returns(s, a)$

$Q(s, a) \leftarrow \text{average}(Returns(s, a))$

(c) For each  $s$  in the episode:

$A^* \leftarrow \arg \max_a Q(s, a)$

For all  $a \in \mathcal{A}(s)$ :

$$\pi(a|s) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon/|\mathcal{A}(s)| & \text{if } a = A^* \\ \varepsilon/|\mathcal{A}(s)| & \text{if } a \neq A^* \end{cases}$$

# Off-Policy Methods

Like exploring starts, soft policies ensure all  $(s, a)$  are visited infinitely many times

- But policies restricted to be soft  
⇒ Optimal policy is usually deterministic!
- Could slowly reduce  $\epsilon$ , but not clear how fast

Other approach: **off-policy learning**

- Learn  $q_\pi$  based on experience generated with *behaviour policy*  $\mu \neq \pi$
- Requires “coverage”: if  $\pi(a|s) > 0$  then  $\mu(a|s) > 0$ , for all  $s, a$   
— e.g. use soft policy  $\mu$
- $\pi$  can be deterministic

## Discussion: On-Policy vs Off-Policy Methods

### On-policy:

Learn  $q_\pi$  and improve  $\pi$  while following  $\pi$

### Off-policy:

Learn  $q_\pi$  and improve  $\pi$  while following  $\mu$

# Importance Sampling Ratio

We have episodes generated from  $\mu$

$\Rightarrow$  Expected return at  $t$  is  $\mathbb{E}_{\mu}[G_t|S_t = s] = v_{\mu}(s)$

Fix expectation with **sampling importance ratio**:

$$\rho_{t:T} \doteq \frac{\prod_{k=t}^{T-1} \pi(A_k|S_k) p(S_{k+1}, R_{k+1}|S_k, A_k)}{\prod_{k=t}^{T-1} \mu(A_k|S_k) p(S_{k+1}, R_{k+1}|S_k, A_k)} = \prod_{k=t}^{T-1} \frac{\pi(A_k|S_k)}{\mu(A_k|S_k)}$$

- $\mathbb{E}_{\mu}[\rho_{t:T} G_t|S_t = s] = v_{\pi}(s)$

# Importance Sampling Ratio

$$\begin{aligned}\mathbb{E}_{\mu}[\rho_{t:T} G_t | S_t = s] &= \sum_{E: S_t = s} \left[ \prod_{k=t}^{T-1} \mu(A_k | S_k) p(S_{k+1}, R_{k+1} | S_k, A_k) \right] \rho_{t:T} G_t \\&= \sum_{E: S_t = s} \left[ \prod_{k=t}^{T-1} \mu(A_k | S_k) p(S_{k+1}, R_{k+1} | S_k, A_k) \right] \prod_{k=t}^{T-1} \frac{\pi(A_k | S_k)}{\mu(A_k | S_k)} G_t \\&= \sum_{E: S_t = s} \left[ \prod_{k=t}^{T-1} \pi(A_k | S_k) p(S_{k+1}, R_{k+1} | S_k, A_k) \right] G_t \\&= v_{\pi}(s)\end{aligned}$$



# Evaluating Policies with Importance Sampling

Denote episodes  $E^i = \langle S_0^i, A_0^i, R_1^i, S_1^i, A_1^i, R_2^i, \dots, S_{T_i}^i \rangle$

Define  $\mathcal{E}(s)/\mathcal{E}(s, a)$  as before for first-visit or every-visit MC

Estimate  $v_\pi/q_\pi$  as

$$v_\pi(s) \approx \eta^{-1} \sum_{t_i \in \mathcal{E}(s)} \rho_{t_i:T_i} G_{t_i}^i$$

$$q_\pi(s, a) \approx \eta^{-1} \sum_{t_i \in \mathcal{E}(s, a)} \rho_{t_i+1:T_i} G_{t_i}^i$$

- **Ordinary** importance sampling:  $\eta = |\mathcal{E}(s, a)|$
- **Weighted** importance sampling:  $\eta = \sum_{t_i \in \mathcal{E}(s)} \rho_{t_i:T_i}$  resp.  $\eta = \sum_{t_i \in \mathcal{E}(s, a)} \rho_{t_i+1:T_i}$

# Off-Policy Value Estimation in Blackjack Example


$\pi$  : stick if player sum is 20 or 21, else hit

$\mu$  : uniformly random


$s$  : player sum 13  
dealer showing 2  
usable ace

True value:

$$v_{\pi}(s) \approx -0.27726$$



# Infinite Variance in Ordinary Importance Sampling



# Iterative Implementation of Weighted Importance Sampling

Input: an arbitrary target policy  $\pi$

Initialize, for all  $s \in \mathcal{S}$ ,  $a \in \mathcal{A}(s)$ :

$Q(s, a) \in \mathbb{R}$  (arbitrarily)

$C(s, a) \leftarrow 0$

Loop forever (for each episode):

$b \leftarrow$  any policy with coverage of  $\pi$

Generate an episode following  $b$ :  $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

$W \leftarrow 1$

Loop for each step of episode,  $t = T-1, T-2, \dots, 0$ , while  $W \neq 0$ :

$G \leftarrow \gamma G + R_{t+1}$

$C(S_t, A_t) \leftarrow C(S_t, A_t) + W$

$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \frac{W}{C(S_t, A_t)} [G - Q(S_t, A_t)]$

$W \leftarrow W \frac{\pi(A_t|S_t)}{b(A_t|S_t)}$

# Off-Policy Monte Carlo Control

Initialize, for all  $s \in \mathcal{S}$ ,  $a \in \mathcal{A}(s)$ :

$Q(s, a) \in \mathbb{R}$  (arbitrarily)

$C(s, a) \leftarrow 0$

$\pi(s) \leftarrow \operatorname{argmax}_a Q(s, a)$  (with ties broken consistently)

Loop forever (for each episode):

$b \leftarrow$  any soft policy

Generate an episode using  $b$ :  $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

$W \leftarrow 1$

Loop for each step of episode,  $t = T-1, T-2, \dots, 0$ :

$G \leftarrow \gamma G + R_{t+1}$

$C(S_t, A_t) \leftarrow C(S_t, A_t) + W$

$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \frac{W}{C(S_t, A_t)} [G - Q(S_t, A_t)]$

$\pi(S_t) \leftarrow \operatorname{argmax}_a Q(S_t, a)$  (with ties broken consistently)

Why?

Hint:  $\pi$  deterministic

Hint:  $q_\pi(s, a)$  uses  $\rho_{t+1:T}$

$\left\{ \begin{array}{l} \text{If } A_t \neq \pi(S_t) \text{ then exit inner Loop (proceed to next episode)} \\ W \leftarrow W \frac{1}{b(A_t|S_t)} \end{array} \right.$

Required:

- RL book, chapter 5 (5.1–5.7)

Optional:

- *Sequential Monte Carlo Methods in Practice*  
Arnaud Doucet, Nando de Freitas, Neil Gordon (editors)  
University library has copies