

# Reinforcement Learning

## Planning and Learning

---

Stefano Albrecht, Pavlos Andreadis


4 February 2020



THE UNIVERSITY of EDINBURGH  
**informatics**


- Planning in reinforcement learning
- Dyna-Q
- Rollout planning
- Monte Carlo tree search
- Offline vs online planning

## Unified View



# Planning

**Planning:** any process which uses a model of the environment to compute a plan of action (policy) to achieve a specified goal



- Dynamic programming is planning: uses model  $p(s', r|s, a)$

# Model

**Model:** anything the agent can use to predict how the environment will respond to its actions

- **Distribution model:** description of all possibilities and their probabilities


$$p(s', r | s, a) \quad \text{for all } s, a, s', r$$

- **Simulation (sample) model:** produces sample outcomes

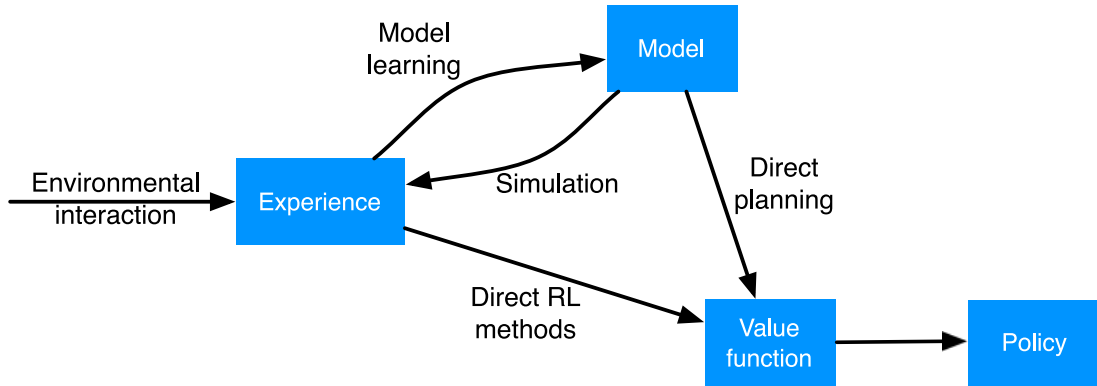
$$\hat{p}(s, a) \rightarrow (\mathcal{S}, \mathcal{R}) \quad \text{s.t.} \quad \Pr\{\hat{p}(s, a) = (s', r)\} = p(s', r | s, a)$$

- Simulation model usually easier to specify than distribution model

## Model-free RL



## Model-based RL



## Dyna-Q: Integrating Planning, Learning, Acting

Initialize  $Q(s, a)$  and  $Model(s, a)$  for all  $s \in \mathcal{S}$  and  $a \in \mathcal{A}(s)$

Do forever:

(a)  $S \leftarrow$  current (nonterminal) state

(b)  $A \leftarrow \varepsilon$ -greedy( $S, Q$ )

(c) Execute action  $A$ ; observe resultant reward,  $R$ , and state,  $S'$

(d)  $Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$  ← **direct RL**

(e)  $Model(S, A) \leftarrow R, S'$  (assuming deterministic environment) ← **model learning**

(f) Repeat  $n$  times:

$S \leftarrow$  random previously observed state



$A \leftarrow$  random action previously taken in  $S$

$R, S' \leftarrow Model(S, A)$

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$  | ← **planning**




# Dyna-Q in Maze Example




# Dyna-Q in Maze Example

Greedy policy halfway through second episode:


WITHOUT PLANNING ( $n=0$ )




WITH PLANNING ( $n=50$ )



## When the Model is Wrong: Blocking Maze



## When the Model is Wrong: Shortcut Maze



Dyna-Q+ uses an exploration bonus heuristic:

- Keeps track of time since each state-action pair was tried in real environment
- Bonus reward is added for transitions caused by state-action pairs related to how long ago they were tried:

$$R + \kappa\sqrt{\tau}$$

time since last visiting  
the state-action pair



- Incentive to re-visit “old” state-action pairs

Dyna-Q uses model to reuse *past* experiences

## Rollout planning:

- Use model to simulate (“rollout”) *future* trajectories
- Each trajectory starts at current state  $S_t$
- Focus usually on finding best action  $A_t$  for state  $S_t$

# Rollout Planning with Forward Updating

## Rollout Q-planning with forward updating:

- 1: Given: simulation model  $Model$
- 2: Initialise:  $Q(s, a)$  for all  $s, a$
- 3: **for**  $t = 0, 1, 2, 3, \dots$  **do**
- 4:    $S_t \leftarrow$  current state
- 5:   **for**  $n$  times ( $n$  rollouts) **do**
- 6:      $S \leftarrow S_t$
- 7:     **while**  $S$  is non-terminal (or fixed-length rollouts) **do**
- 8:       select action  $A$  based on  $Q(S, \cdot)$  with some exploration // e.g.  $\epsilon$ -greedy
- 9:        $(R, S') \sim Model(S, A)$
- 10:       Q-update:  $Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$
- 11:        $S \leftarrow S'$
- 12:   select action  $A_t$  greedily from  $Q(S_t, \cdot)$

# Rollout Planning

If model **correct** and under Q-learning conditions (all  $(s, a)$  infinitely visited and standard  $\alpha$ -reduction), rollout planning learns optimal policy

If model **incorrect**, learned policy likely sub-optimal on real task

- Can range from slightly sub-optimal to failing to solve real task (examples?)

Next: can we use rewards from rollouts more effectively?

⇒ Back-propagate rewards




# Rollout Planning with Backward Updating (Back-Propagation)

## Rollout Q-planning with backward updating:

- 1: Given: simulation model *Model*
- 2: Initialise:  $Q(s, a)$  for all  $s, a$ ; LIFO stack  $Trace = \{\}$
- 3: **for**  $t = 0, 1, 2, 3, \dots$  **do**
- 4:    $S_t \leftarrow$  current state
- 5:   **for**  $n$  times ( $n$  rollouts) **do**
- 6:      $S \leftarrow S_t$
- 7:     **while**  $S$  is non-terminal (or fixed-length rollouts) **do**    *// Rollout*
- 8:       select action  $A$  based on  $Q(S, \cdot)$  with some exploration
- 9:        $(R, S') \sim Model(S, A)$
- 10:       push  $(S, A, R, S')$  to  $Trace$
- 11:        $S \leftarrow S'$
- 12:     **while**  $Trace$  not empty **do**                                   *// Backprop*
- 13:       pop  $(S, A, R, S')$  from  $Trace$
- 14:        $Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$
- 15:     select action  $A_t$  greedily from  $Q(S_t, \cdot)$


# Rollout Planners in Maze Example



$$\gamma = 0.95$$

$$\epsilon = 0.1$$

$$\alpha = 0.1$$




# Monte Carlo Tree Search

## Monte Carlo Tree Search (MCTS):


- General, efficient rollout planner
- Stores **partial**  $Q$  as tree and **asymmetrically expands** tree based on most promising actions

$Q$  is recursive tree structure:

$$Q(s, a) = \mathbb{E}[R_{t+1} + \gamma \max_{a'} Q(S_{t+1}, a') \mid S_t = s, A_t = a]$$



# Phases of Monte Carlo Tree Search



Browne et al. (2012)

## MCTS-Search( $S_t$ ):

- 1: Find node  $v_0$  with  $state(v_0) = S_t$  (or create new node)
- 2: **while** within computational budget **do**
- 3:    $v_l \leftarrow TreePolicy(v_0)$
- 4:    $\Delta \leftarrow DefaultPolicy(state(v_l))$
- 5:    $Backprop(v_l, \Delta)$
- 6: **return**  $action(BestChild(v_0))$    *// e.g. most visited child; highest expected return*

- Backprop works just as before
- Tree policy can be any exploration policy (balancing exploration and exploitation)

# Upper Confidence Bounds for Trees

## Upper Confidence Bounds for Trees (UCT):

- Uses UCB action selection as tree policy and  $\alpha = 1/N(S, A)$
- Popular MCTS variant: easy to use and effective

UCB recap: estimate **upper bound** on action value:

$$A \leftarrow \begin{cases} a \text{ if not tried before in } S \\ \arg \max_a Q(S, a) + c \sqrt{\log N(S) / N(S, a)} \end{cases}$$

- $N(S)$  is number of times state  $S$  has been visited
- $N(S, a)$  is number of times action  $a$  was selected in  $S$

# Simulation Step

Simulation step gives estimate of return at state, e.g.:

## Random-DefaultPolicy( $S$ ):

```
1:  $G \leftarrow 0$ 
2: while  $S$  is non-terminal do
3:    $A \leftarrow$  random action (uniformly)
4:    $(R, S') \sim \text{Model}(S, A)$ 
5:    $G \leftarrow R + \gamma G$ 
6:    $S \leftarrow S'$ 
7: return  $G$ 
```

Possible improvements:

- Average over multiple simulations
- Use domain-specific heuristic to
  - select better actions than random
  - evaluate state directly (e.g. in Chess we know that some states are better than others)

# Offline Planning

Imagine you are given an MDP for a chess game against a specific opponent:

## Offline planning:

- Use MDP to find best policy **before** the actual chess game takes place (offline)
- Use as much time as needed to find policy
- Policy is **complete**: gives optimal action for all possible states

Dyna-Q and dynamic programming are suitable for offline planning





# Online Planning

Imagine you are given an MDP for a chess game against a specific opponent:

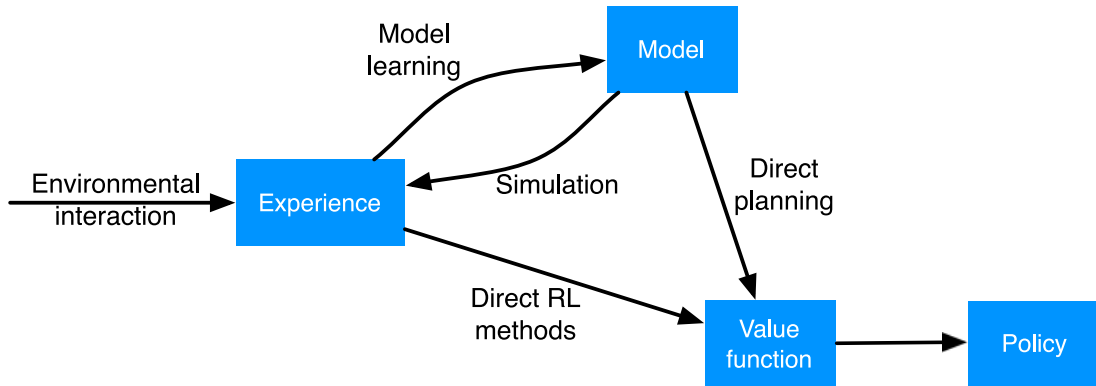
## Online planning:

- Use MDP to find best policy **during** the actual chess game (online)
- Limited computing time budget at each state (e.g. seconds/minutes in chess)
- Policy usually **incomplete**: gives optimal action for *current state*

Rollout planning (including MCTS) are suitable for online planning



## Paths to a Policy: Model-Based RL



Required:

- RL book, chapter 8 (8.1–8.6, 8.8–8.11)

Optional:

- Browne et al. (2012). A Survey of Monte Carlo Tree Search Methods. IEEE Transactions on Computational Intelligence and AI in Games, Vol. 4, No. 1
- UCT paper: L. Kocsis and C. Szepesvari (2006). Bandit based Monte-Carlo Planning. European Conference on Machine Learning
- T. Vodopivec, S. Samothrakis, B. Ster (2017). On Monte Carlo Tree Search and Reinforcement Learning. Journal of Artificial Intelligence Research, Vol. 60